

Fig 5 — Training Stability: Policy Entropy and KL Divergence

